

R: Split/Apply/Combine using the Orange dataset

```
# clear workspace  
rm(list=ls())  
# Get info on "Orange" R dataset  
help("Orange")  
# Make a copy of the "Orange" template and save as my  
orange <- as.data.frame(Orange)
```

"If I had an hour to solve a problem, I'd spend 55 minutes thinking about the problem and 5 minutes thinking about solutions."

— [Albert Einstein](#)


R Orange dataset = original template Growth of Orange Trees

Description

The Orange data frame has 35 rows and 3 columns of records of the growth of orange trees.

Columns:

- **Tree** - ordered factor indicating the tree on which the measurement is made.
- **age** - numeric vector giving the age of the tree (days since 1968/12/31)
- **circumference** - numeric vector of trunk circumferences (mm).



```
> str(orange) # Get structure
'data.frame':   35 obs. of  3 variables:
 $ Tree      : Ord.factor w/ 5 levels "3"<"1"<"5"<"2"<...:
 $ age       : num  118 484 664 1004 1231 ...
 $ circumference: num  30 58 87 115 120 142 145 33 69 111 ...
```

What are values of 'Tree' factor?

```
> levels(orange$Tree)
[1] "3" "1" "5" "2" "4"
```

7 "ages" or data collection dates:

	Tree.age	Tree.date
1	118	1969-04-27
2	484	1970-04-28
3	664	1970-10-25
4	1004	1971-09-30
5	1231	1972-05-14
6	1372	1972-10-02
7	1582	1973-04-30

Data Exploration: "orange" data frame

```
# R/Studio Viewer
> view(orange)
```

Data for Tree 1

Data for Tree 2

Data for Tree 3

Data for Tree 5

35 observations of 3 variables

	Tree	age	circumference
1	1	118	30
2	1	484	58
3	1	664	87
4	1	1004	115
5	1	1231	120
6	1	1372	142
7	1	1582	145
8	2	118	33
9	2	484	69
10	2	664	111
11	2	1004	156
12	2	1231	172
13	2	1372	203
14	2	1582	203
15	3	118	30
16	3	484	51
17	3	664	75

Each Tree has 7 "ages" or dates when data was collected.

```

attach(orange) # Can now skip orange$ prefix
# What are min/max values of 'age'?
# Get X-axis range for: age
Xrange <- range(age); Xrange


[1] 118 1582

# Get Y-axis range for: circumference
Yrange <- range(circumference); Yrange

[1] 30 214

# See help("plot"), example("plot")
plot(age, circumference,
      main="Orange Tree Growth",
      xlab="Tree age(days since 1968/12/31)",
      ylab="Tree circumference (mm)",
      xlim= Xrange, ylim= Yrange);

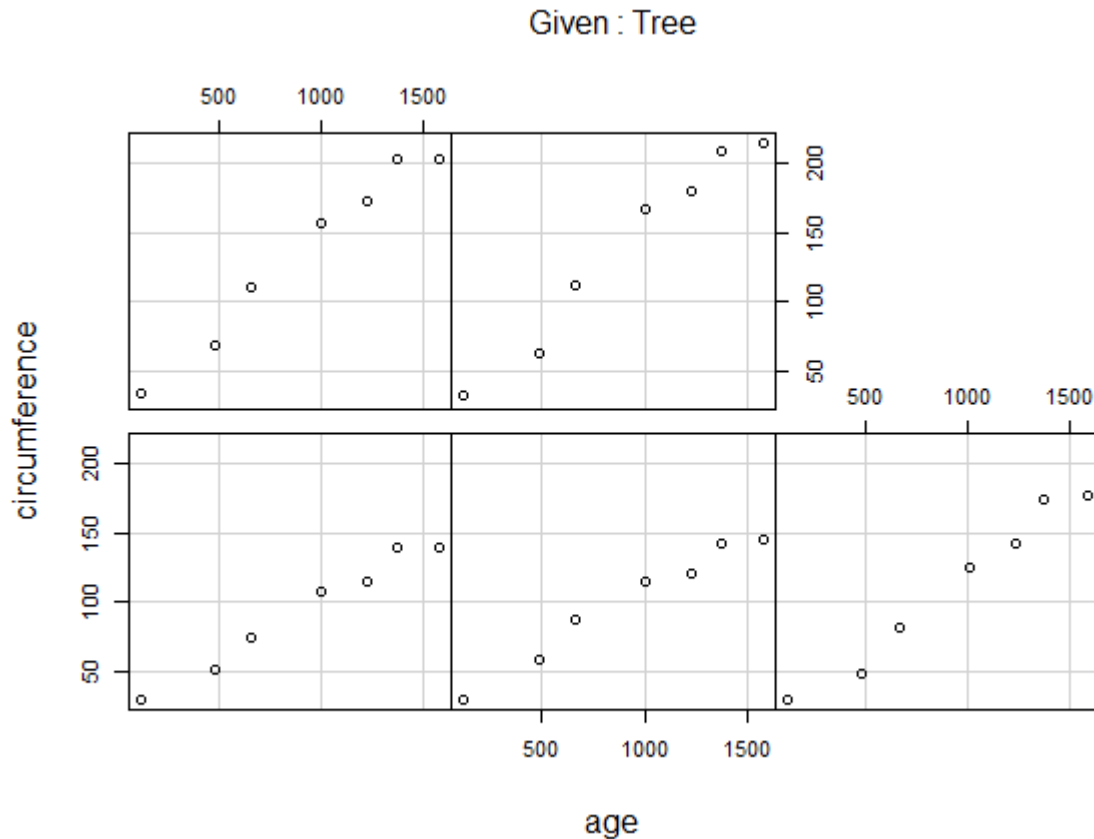
```



What does this mean?

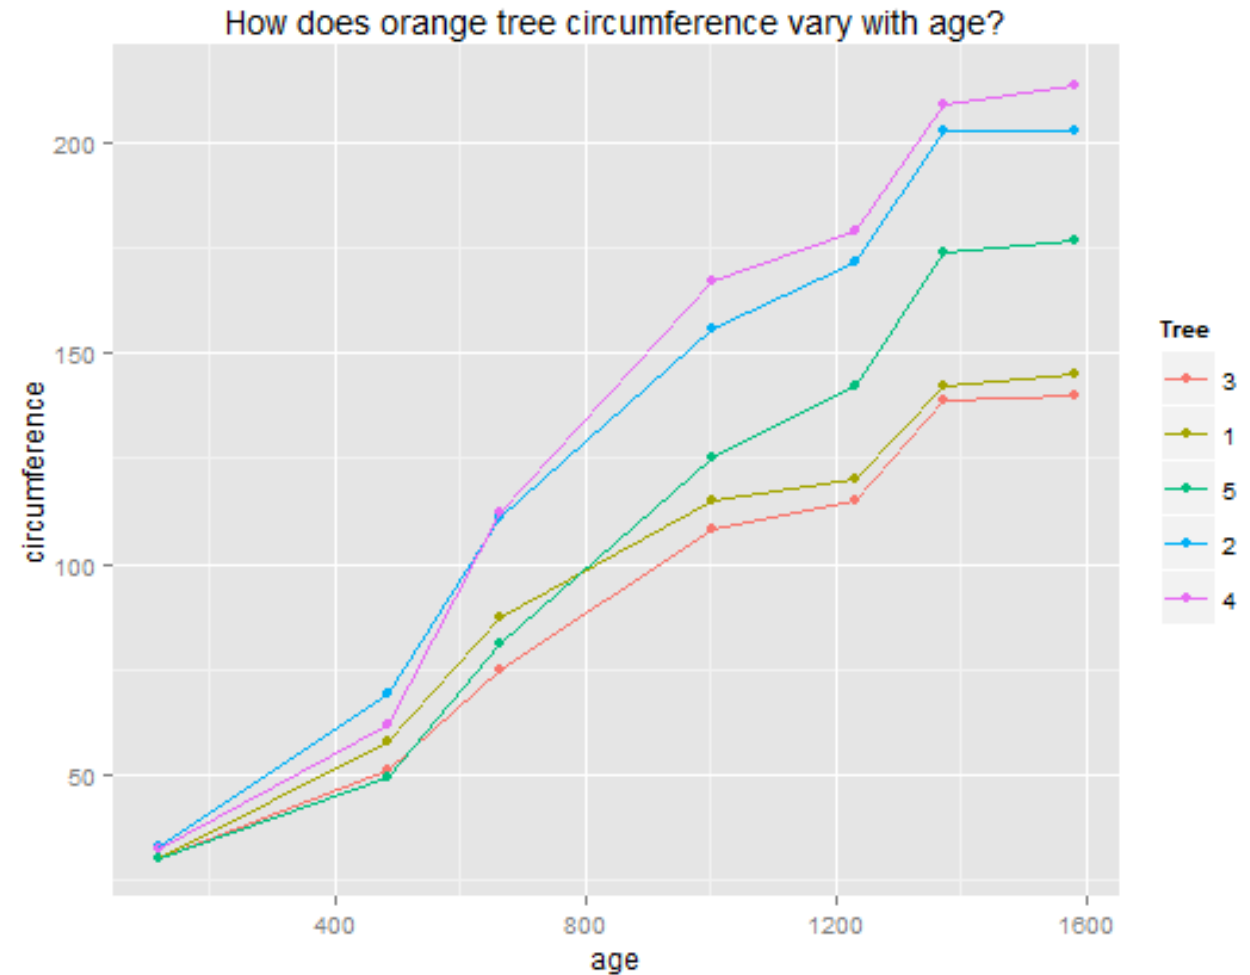
- Data was collected for each Tree at 7 "ages" : 118, 484, 664, 1004, 1231, 1372, 1582.
- Each dot is the circumference measurement for a Tree.
- As age increases, circumference also increases.

Other plots of the same data...



From `help("orange")`

```
# scatter plots for each Tree
require(stats); require(graphics)
coplot(circumference ~ age | Tree,
       data = Orange, show.given = FALSE)
```



From [ggplot2-tutorial](#)

```
library(ggplot2)
qplot(age, circumference, data = Orange,
      geom = c("point", "line"), colour = Tree,
      main = "How does orange tree circumference vary with age?")
```

```
# Get rows from orange where Tree = "1"
```

```
# Method 1: Use subset()
```

```
Tree1 <- subset(orange, orange$Tree == "1")
```

```
# Method 2: Use sqldf()
```

```
library(sqldf)
```

```
Tree1 <- sqldf("SELECT * FROM orange  
              WHERE Tree='1'")
```

```
# Compute correlation. See help("cor")
```

```
cor(Tree1$age, Tree1$circumference)
```

```
[1] 0.9854675
```


What does this mean?

There is a HIGH POSITIVE correlation between age and circumference.

Method 1 & 2 result in the same data frame:

```
> Tree1  
  Tree age circumference  
1    1  118            30  
2    1  484            58  
3    1  664            87  
4    1 1004           115  
5    1 1231           120  
6    1 1372           142  
7    1 1582           145
```

Tree1 fitted model



Problem: Define Covariance/Correlation Chart

Row	Tree	Covariance	Correlation
1	1	<code>cov(age, circumference) where Tree="1"</code>	<code>cor(age, circumference) where Tree="1"</code>
2	2	<code>cov(age, circumference) where Tree="2"</code>	<code>cor(age, circumference) where Tree="2"</code>
3	3	<code>cov(age, circumference) where Tree="3"</code>	<code>cor(age, circumference) where Tree="3"</code>
4	4	<code>cov(age, circumference) where Tree="4"</code>	<code>cor(age, circumference) where Tree="4"</code>
5	5	<code>cov(age, circumference) where Tree="5"</code>	<code>cor(age, circumference) where Tree="5"</code>

*“A problem well defined
is a problem half-solved.”*
– [John Dewey](#)

We'll try 3 methods:

- *Using `subset()` and `merge()`
- *Using `tapply()`
- Using `ddply()`

*Adapted from MyFirstRLesson.R
Patricia Hoffman, PhD.

<http://patriciahoffmanphd.com/startr.php>

UCSC Extension – Course 262

Data Mining and Machine Learning

SubProblems:

- Get **chart1** for orange
- Get **chart2** for oPlus10 (circumference +10)
- Get **chart3** for oX2 (circumference *2)
- Get **chart4** for oX_2 (circumference * -2)

How to do all of the above without cutting/pasting code for the same functionality but different inputs?

Solution:

Define & re-use function **covcor()**.

```
chart1 <- covcor(orange)
```

```
oPlus10 <- orange  
oPlus10$circumference <- oPlus10$circumference+10  
chart2 <- covcor(oPlus10)
```

```
oX2 <- orange  
oX2$circumference <- oX2$circumference*(-2)  
chart3 <- covcor(oX2)
```

```
oX_2 <- orange  
oX_2$circumference <- oX_2$circumference*(-2)  
chart4 <- covcor(oX_2)
```

Method 1: Define function `covcor()`

Using `subset()` and `merge()`

```
# FUNCTION covcor:  
# INPUT: data frame with Tree, age, circumference  
# OUTPUT: data frame with Tree,  
Correlation/Covariance
```

```
covcor <- function(orangedf)  
{  
  # SPLIT using subset()  
  Tree1 <- subset(orangedf, Tree == "1")  
  Tree2 <- subset(orangedf, Tree == "2")  
  Tree3 <- subset(orangedf, Tree == "3")  
  Tree4 <- subset(orangedf, Tree == "4")  
  Tree5 <- subset(orangedf, Tree == "5")  
  
  # APPLY functions cov() and cor()  
  # Calculate COVARIANCE for each Tree  
  cov.df <- data.frame("fac" = factor(1:5),  
    "cov" = c(cov(Tree1$age,Tree1$circumference),  
    cov(Tree2$age,Tree2$circumference),  
    cov(Tree3$age,Tree3$circumference),  
    cov(Tree4$age,Tree4$circumference),  
    cov(Tree5$age,Tree5$circumference)))
```

```
  # Calculate CORRELATION for each Tree  
  cor.df <- data.frame("fac" = factor(1:5),  
    "cor" = c(cor(Tree1$age,Tree1$circumference),  
    cor(Tree2$age,Tree2$circumference),  
    cor(Tree3$age,Tree3$circumference),  
    cor(Tree4$age,Tree4$circumference),  
    cor(Tree5$age,Tree5$circumference)))  
  # COMBINE using merge()  
  # "Join" cov.df and cor.df using the fac column  
  chart <- merge(cov.df, cor.df,by.x = "fac", by.y =  
    "fac")  
  # Assign column names  
  names(chart) <- c("TREE","COVARIANCE","CORRELATION")  
  # Assign row names  
  row.names(chart)<- c("1","2","3","4","5")  
  return(chart)  
}
```


INPUT: data frame orangedf

Step 1: SPLIT

```
Tree1 <- subset(orangedf, Tree == "1")
Tree2 <- subset(orangedf, Tree == "2")
Tree3 <- subset(orangedf, Tree == "3")
Tree4 <- subset(orangedf, Tree == "4")
Tree5 <- subset(orangedf, Tree == "5")
```

1	118	30
1	484	58
1	664	87

Tree1

1	2	118	33
1	2	484	69
1	2	664	111

Tree2

1	2	3	118	30
1	2	3	484	51
1	2	3	664	75

Tree3

1	2	3	4	118	32
1	2	3	4	484	62
1	2	3	4	664	112

Tree4

1	2	3	4	5	118	30
1	2	3	4	5	484	49
1	2	3	4	5	664	81
1	2	3	4	5	1004	125
1	2	3	4	5	1231	142
1	2	3	4	5	1372	174
1	2	3	4	5	1582	177

Tree5

```
# Calculate the COVARIANCE for each tree
cov.df <- data.frame("fac" = factor(1:5),
"cov" = c(cov(Tree1$age,Tree1$circumference),
cov(Tree2$age,Tree2$circumference),
cov(Tree3$age,Tree3$circumference),
cov(Tree4$age,Tree4$circumference),
cov(Tree5$age,Tree5$circumference)))
```

Step 2: APPLY

cov.df

	fac	V2
1	1	cov(age,circumference)
2	2	cov(age,circumference)
3	3	cov(age,circumference)
4	4	cov(age,circumference)
5	5	cov(age,circumference)

```
# Calculate the CORRELATION for each Tree
cor.df <- data.frame("fac" = factor(1:5),
"cor" = c(cor(Tree1$age,Tree1$circumference),
cor(Tree2$age,Tree2$circumference),
cor(Tree3$age,Tree3$circumference),
cor(Tree4$age,Tree4$circumference),
cor(Tree5$age,Tree5$circumference)))
```

cor.df

	fac	V2
1	1	cor(age,circumference)
2	2	cor(age,circumference)
3	3	cor(age,circumference)
4	4	cor(age,circumference)
5	5	cor(age,circumference)

Step 3: COMBINE

```
chart <- merge(cov.df, cor.df,by.x = "fac", by.y = "fac")
names(chart) <- c("TREE","COVARIANCE","CORRELATION")
row.names(chart)<- c("1","2","3","4","5")
return(chart)
```

chart

	TREE	COVARIANCE	CORRELATION
1	1	22340.07	0.9854675
2	2	34290.45	0.9873624
3	3	22239.83	0.9881766
4	4	37062.62	0.9844610
5	5	30442.81	0.9877376

OUTPUT:
data frame
chart

Method 2: Define function covcor()

Using **tapply(index(array), column, function)+ merge()**

```
covcor <- function(orange)
{
# SPLIT/APPLY
agg.cor <- tapply(1:nrow(orange),
                 orange$Tree, FUN=function(x)
                 cor(orange$age[x],
                     orange$circumference[x]))
agg.cov <- tapply(1:nrow(orange),
                 orange$Tree, FUN=function(x)
                 cov(orange$age[x],
                     orange$circumference[x]))
# Convert arrays to tables, then dfs
agg.cor <- (as.data.frame(as.table(agg.cor)))
agg.cov <- (as.data.frame(as.table(agg.cov)))
# COMBINE
chart <- merge(agg.cov, agg.cor, by=c("Var1"))
colnames(chart) <- c("TREE", "COVARIANCE",
                    "CORRELATION")

return(chart)
}
```

- tapply applies **function** to elements in **index(array)** and groups results by **column**.
- Outputs an array, which must be converted to table, then data frame for merge().
- Similar to SQL "Group By" concept:

```
CREATE VIEW aggcov AS (SELECT Tree,
*COVAR_POP(age,circumference) AS COVARIANCE FROM
orange GROUP BY Tree);
```

```
CREATE VIEW aggcor AS (SELECT Tree,
*CORRELATION(age,circumference) AS CORRELATION
FROM orange GROUP BY Tree);
```

```
CREATE VIEW chart as (SELECT aggcov.Tree as TREE,
COVARIANCE, CORRELATION FROM aggcov, aggcor WHERE
aggcov.Tree = aggcor.Tree);
```

*Currently no built-in CORRELATION functions in SQL.
See: <http://www.xarg.org/2012/07/statistical-functions-in-mysql/>

INPUT:
data
frame
orangedf

Step 1: SPLIT/APPLY

```
agg.cov <-tapply(1:nrow(orangedf),  
orangedf$Tree, FUN=function(x)  
cov(orangedf$age[x],  
orangedf$circumference[x]))
```

Values	
agg. cor	array[5]
agg. cov	array[5]

```
agg.cor <-tapply(1:nrow(orangedf),  
orangedf$Tree, FUN=function(x)  
cor(orangedf$age[x],  
orangedf$circumference[x]))
```

```
agg.cov <-  
(as.data.frame  
(as.table(agg.cov)))
```

	Var1	Freq
1	3	22239.83
2	1	22340.07
3	5	30442.81
4	2	34290.45
5	4	37062.62

	Var1	Freq
1	3	0.9881766
2	1	0.9854675
3	5	0.9877376
4	2	0.9873624
5	4	0.9844610

```
agg.cor <-  
(as.data.frame  
(as.table(agg.cor)))
```

Step 2: COMBINE

```
chart <-merge(agg.cov, agg.cor, by=c("Var1"))  
colnames(chart) <- c("TREE","COVARIANCE","CORRELATION")
```

chart

	TREE	COVARIANCE	CORRELATION
1	1	22340.07	0.9854675
2	2	34290.45	0.9873624
3	3	22239.83	0.9881766
4	4	37062.62	0.9844610
5	5	30442.81	0.9877376

OUTPUT:
data frame
chart

Method 3: Define function `covcor()`

Using `ddply(data frame, column, function1, function2)`

```
library(plyr)
# Define 2 new functions to be used with ddply
ncor <- function(newdf) {chart <- cor(newdf[, 2],
  newdf[, 3]); return(chart)}
ncov <- function(newdf) {chart <- cov(newdf[, 2],
  newdf[, 3]);return(chart)}

# Now use ddply to compute BOTH ncov and ncor:
covcor <- function(odf) {
  chart <- ddply(odf, .(Tree), c("ncov", "ncor"))
  colnames(chart) <- c("TREE", "COVARIANCE", "CORRELATION")
  return(chart) }

# That's it! no subsetting or merging required!
orange <- as.data.frame(Orange)
o.chart <- covcor(Orange)
```

From help("plyr"):

- **ddply**: Split data frame, apply function(s), and return results in a data frame.
- **plyr** is a set of tools that implement the **split-apply-combine (SAC)** pattern.

SAC is essential in data analysis. To solve a problem:

- Break it down into smaller pieces.
- Do something to each piece.
- Combine the results back together again.

INPUT:
data
frame
odf

Step 0: Define new functions

```
library(plyr)
# Define 2 new functions to be used with ddply
ncor <- function(newdf) {chart <- cor(newdf[, 2], newdf[, 3]); return(chart)}
ncov <- function(newdf) {chart <- cov(newdf[, 2], newdf[, 3]);return(chart)}
```

Step 1: SPLIT/APPLY/COMBINE

Now use ddply to compute BOTH ncov and ncor:

```
covcor <- function(odf)
{chart <- ddply(odf, .(Tree), c("ncov", "ncor"))
  colnames(chart) <- c("TREE", "COVARIANCE", "CORRELATION")
  return(chart)}
```

That's it! no subsetting or merging required!

```
orange <- as.data.frame(Orange)
chart <- covcor(orange)
```

chart

	TREE	COVARIANCE	CORRELATION
1	1	22340.07	0.9854675
2	2	34290.45	0.9873624
3	3	22239.83	0.9881766
4	4	37062.62	0.9844610
5	5	30442.81	0.9877376

OUTPUT:
data frame
chart

DATA TRANSFORMATIONS: What happens when we change circumference?



<pre>corcov(orange) TREE COVARIANCE CORRELATION 1 1 22340.07 0.9854675 2 2 34290.45 0.9873624 3 3 22239.83 0.9881766 4 4 37062.62 0.9844610 5 5 30442.81 0.9877376</pre> <p>Original</p>	<pre>corcov(orangeplus10) TREE COVARIANCE CORRELATION 1 1 22340.07 0.9854675 2 2 34290.45 0.9873624 3 3 22239.83 0.9881766 4 4 37062.62 0.9844610 5 5 30442.81 0.9877376</pre> <p>Transformation: circumference+10 Result: No change.</p>
<pre>corcov(orangex2) TREE COVARIANCE CORRELATION 1 1 44680.14 0.9854675 2 2 68580.90 0.9873624 3 3 44479.67 0.9881766 4 4 74125.24 0.9844610 5 5 60885.62 0.9877376</pre> <p>Transformation: circumference*2 Result: Covariance doubles, but correlation stays the same.</p>	<pre>corcov(orangex_2) TREE COVARIANCE CORRELATION 1 1 -44680.14 -0.9854675 2 2 -68580.90 -0.9873624 3 3 -44479.67 -0.9881766 4 4 -74125.24 -0.9844610 5 5 -60885.62 -0.9877376</pre> <p>Transformation: circumference*(-2) Result: Both Covariance & Correlation are negative. Absolute value of Covariance doubles, but absolute value of Correlation stays the same.</p>

Conclusions:

- **Covariance is affected by scale.** Doubling circumference increases covariance (data points "stretch out" to fill more space) but doesn't affect correlation (direction is still positive/upward).
- **Negative correlation/covariance reverses the direction of the relationship.** I.e., as age increases, the tree circumference "shrinks" instead of getting larger.

DATA VISUALIZATION: What happens when we change circumference?

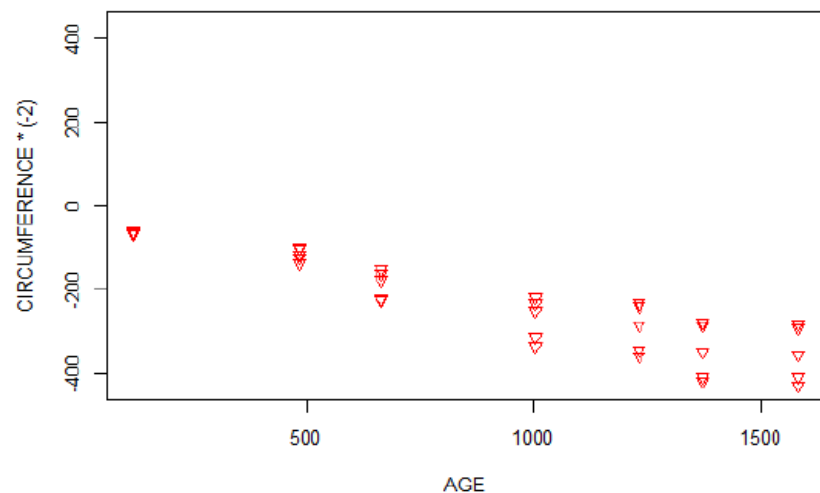
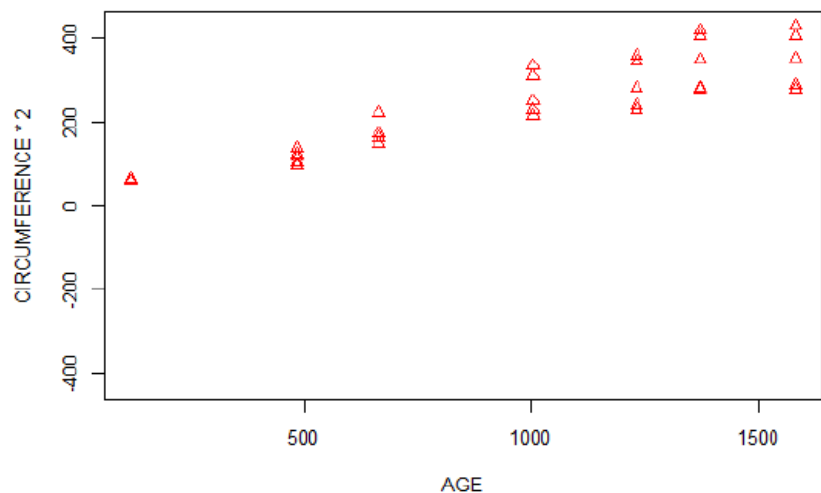
Original



+10:
Minimal change.

X 2:

- **Covariance doubles** (points spread out).
- **Correlation stays the same** (positive relationship.)



X -2:

- **Covariance doubles** (points spread out).
- **Correlation reverses** (negative relationship).

BTW....

7 "ages" or sample dates:

	Tree.age	Tree.date
1	118	1969-04-27
2	484	1970-04-28
3	664	1970-10-25
4	1004	1971-09-30
5	1231	1972-05-14
6	1372	1972-10-02
7	1582	1973-04-30

```
# convert 'age' to a date
secs.per.day <- 24 * 60 * 60
orange$age.sec <- orange$age * secs.per.day
orange$Date <- as.POSIXct(orange$age.sec,
origin = "1968-12-31", tz = "EST")
# Get YYYY-MM-DD
orange$age.date <- substr(orange$Date,1,10)

# Can use getYear, getMonth, getDay,
# getHour, getMin, getSec.
library(gdata)
getYear(orange$Date)
```